

# Design YouTube/Netflix/Hulu

Shawn K.

# YouTube

(Netflix, Hulu)

Design

1. **Intro**
  2. **Requirements**
    - Considerations
    - FR, NFR
  3. **High Level Design**
    - Capacity Estimation
    - Components
    - Video Uploading Workflow
    - Video Streaming Workflow
  4. **Detailed Design**
    - Video Transcoding
    - Optimizations
    - Error Handling
  5. **Wrap Up**
    - Scalability: API, DB
    - Live Streaming
    - Video Takedowns
-

# 1. Intro - YouTube

## 2020 Statistics:

- MAU (Monthly Active Users): 2 billions
- # of videos watched: 5 billions / day
- US adults: 73%
- Video creators: 50 millions
- Ad revenue: \$15.1 billions (2019, up 36% from 2018)
- Mobile traffic: 37%
- Languages: 80 ea.

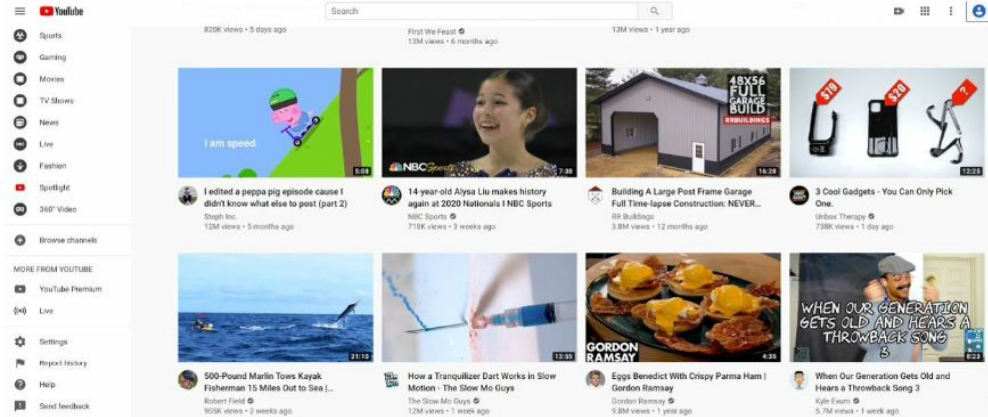
- Content creators upload videos
- Viewers click play.



Simple?



There are lots of complex technologies underneath the simplicity.



## 2. Requirements - Consideration

### General Features

- Watch a video
- Comment
- Share
- Like a video
- Save a video to playlists
- Subscribe to a channel
- etc.

Clarification for 45-60 min



### Understand Problem & Establish Scope

- What features are important? → Ability to upload and watch a video
- What clients do we need to support? → Mobile apps, web browsers, smart TV
- How many daily active users do we have? → 5 million
- What is the average daily time spent on the product? → 30 minutes
- Do we need to support international users? → Yes
- What are the supported video resolutions?
  - The system accepts most of video resolutions & formats
- Is encryption required? → Yes
- Any file size requirement for videos? → max: 1GB
- Can we leverage some of the existing cloud infrastructures provided by GCP, AWS, Azure?
  - Great question! Building everything from scratch is unrealistic for most companies, it is recommended to leverage some of the existing cloud services.



## 2. Requirements - FR

### General Features

- Watch a video
- Comment
- Share
- Like a video
- Save a video to playlists
- Subscribe to a channel
- etc.

Clarification for 45-60 min

### Functional Requirements

- Upload a video and watch a video
- Change video quality
- Clients: mobile apps, web browsers, TV



## 2. Requirements - NFR

### Scalability

- DAU (Daily Active Users): 5 million
- Average Daily Time spent on product: 30 minutes
- International users

### Performance

- **Smooth video streaming:** Users can stream videos in real-time without any lag or latency issue

### High Availability

### Reliability

- The system will be highly reliable

### Security

- Encryption

### Cost

- Low infrastructure cost

# 3. High Level Design - Capacity Estimation (1/2)

[Assumption]

- **Network**

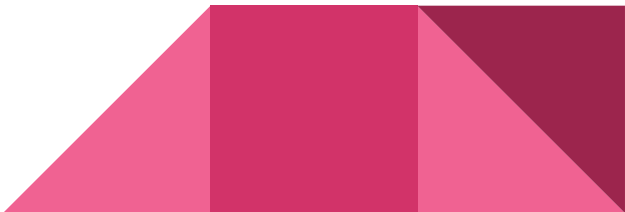
- **DAU** (Daily Active Users): 5 million / day
- **# of videos watched / day**: 5 videos per day
- **# of videos uploaded / day**: 10% of users upload 1 video per day = 5 million (DAU) \* 10% = 500K
- **Average video size**: 300MB

- **Storage**

- **Per Day**: 5 million \* 10% \* 300MB = 500,000 \* 300 = 150,000,000 MB = 150 TB
- **Per Month**: 150TB \* 30 = 4,500 TB = 4.5 PB
- **5 Years**: 4.5 PB \* 12 months \* 5 years = 270 EB

- **CDN Cost**

- CloudFront: 100% of traffic (US), average cost per GB: \$0.02
- 5 million \* 5 videos \* 0.3GB \* \$0.02 = \$150K per day



### 3. High Level Design - Capacity Estimation (2/2)

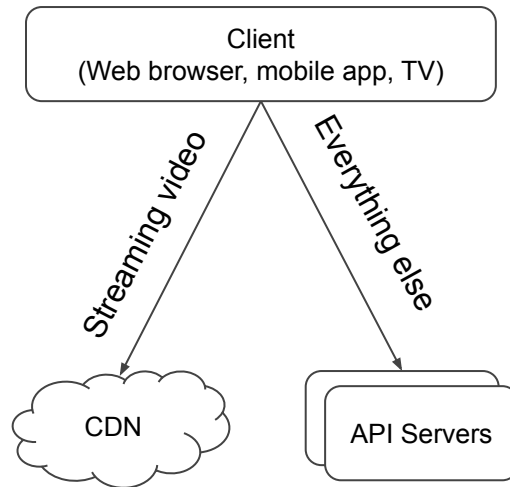
AWS CDN CloudFront: On-demand pricing for Data Transfer to the internet (per GB)

Per Month	United States & Canada	Europe & Israel	South Africa & Middle East	South America	Japan	Australia	Singapore, South Korea, Taiwan, Hong Kong, & Philippines	India
First 10TB	\$0.085	\$0.085	\$0.110	\$0.110	\$0.114	\$0.114	\$0.140	\$0.170
Next 40TB	\$0.080	\$0.080	\$0.105	\$0.105	\$0.089	\$0.098	\$0.135	\$0.130
Next 100TB	\$0.060	\$0.060	\$0.090	\$0.090	\$0.086	\$0.094	\$0.120	\$0.110
Next 350TB	\$0.040	\$0.040	\$0.080	\$0.080	\$0.084	\$0.092	\$0.100	\$0.100
Next 524TB	\$0.030	\$0.030	\$0.060	\$0.060	\$0.080	\$0.090	\$0.080	\$0.100
Next 4PB	\$0.025	\$0.025	\$0.050	\$0.050	\$0.070	\$0.085	\$0.070	\$0.100
Over 5PB	\$0.020	\$0.020	\$0.040	\$0.040	\$0.060	\$0.080	\$0.060	\$0.100



### 3. High Level Design - Components

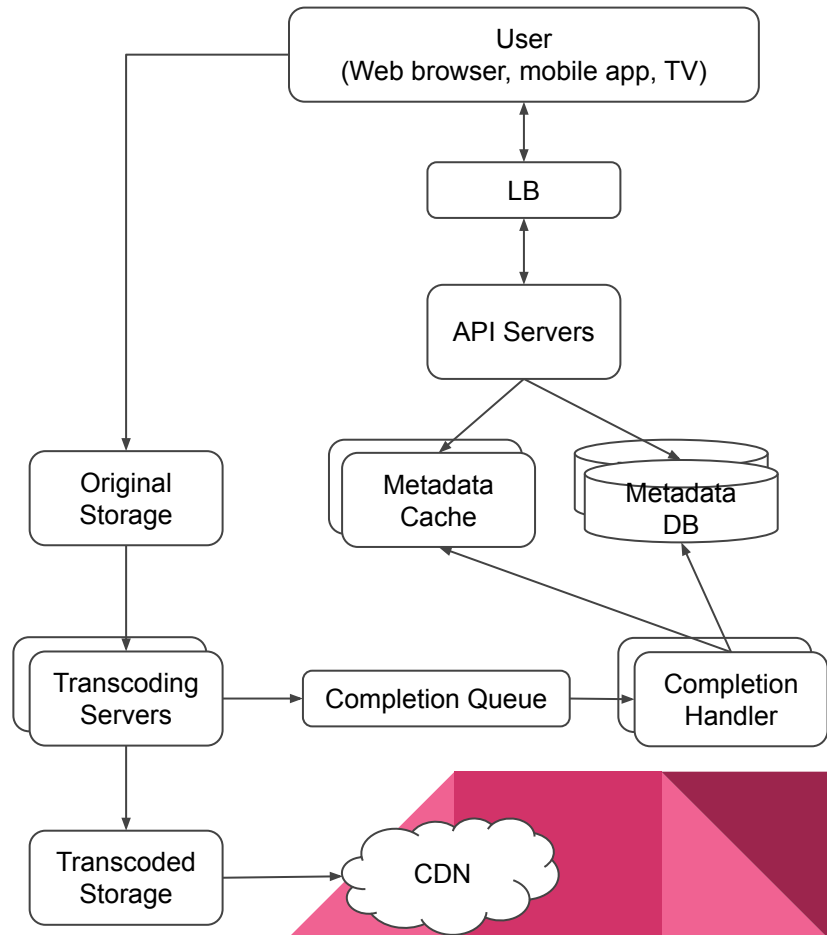
- **Client**
  - Watch YouTube on computer web browser, mobile phone, smart TV
- **CDN**
  - Videos are stored in CDN.
  - When pressing play → a video is streamed from the CDN.
- **API Servers**
  - Everything else except video streaming goes through API servers.
  - Feed recommendation
  - Generating video upload URL
  - Updating metadata database & cache
  - User signup



# 3. High Level Design

## (Video Uploading Workflow)

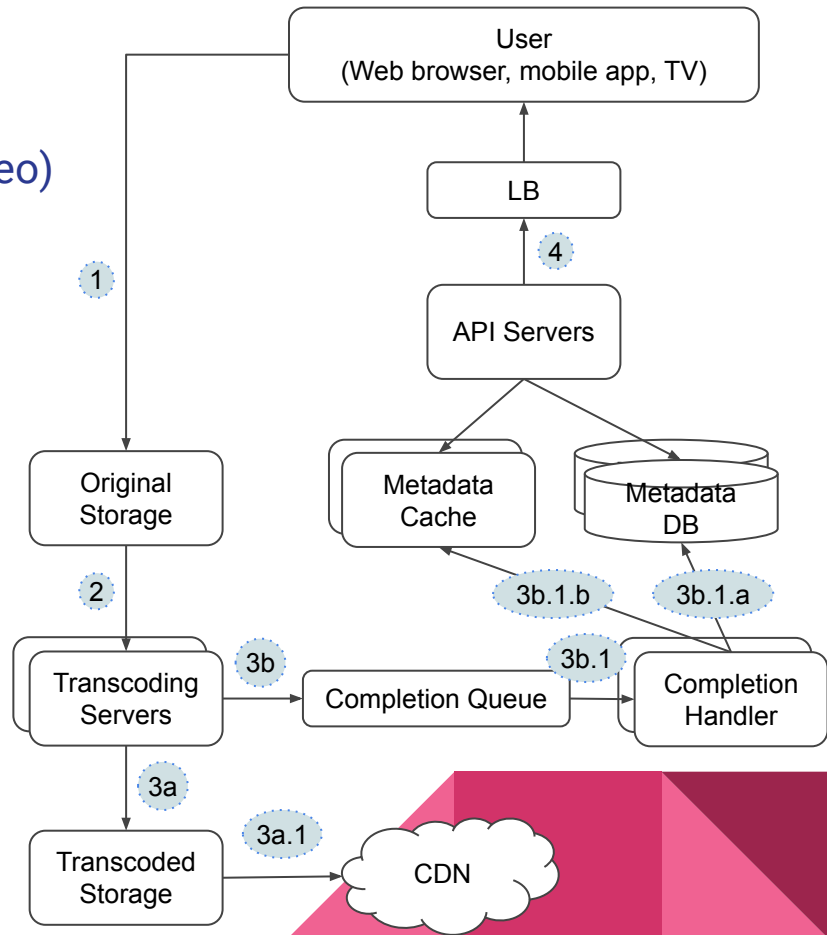
- **User, LB**
- **API Servers:** handles all reqs w/o video streaming
- **Metadata DB**
  - Stores video metadata.
  - Sharding/Replication for high performance/availability
- **Metadata Cache:** for better performance
  - Video metadata
  - User objects
- **Original Storage:** store original videos (BLOB)
- **Transcoding Servers**
  - Video encoding: convert video format
  - Provides best video streams for different devices and bandwidth
- **Transcoding Storage:** store transcoded video files (MPEG, HLS)
- **CDN**
  - Videos are cached in CDN.
  - When clicking 'play' button → a video is streamed from CDN
- **Completion Queue**
  - Message queue: stores info (transcoding completion event)
- **Completion Handler**
  - List of workers: pull event data from completion queue
  - Update metadata cache and database



# 3. High Level Design

(Video Uploading Workflow - upload actual video)

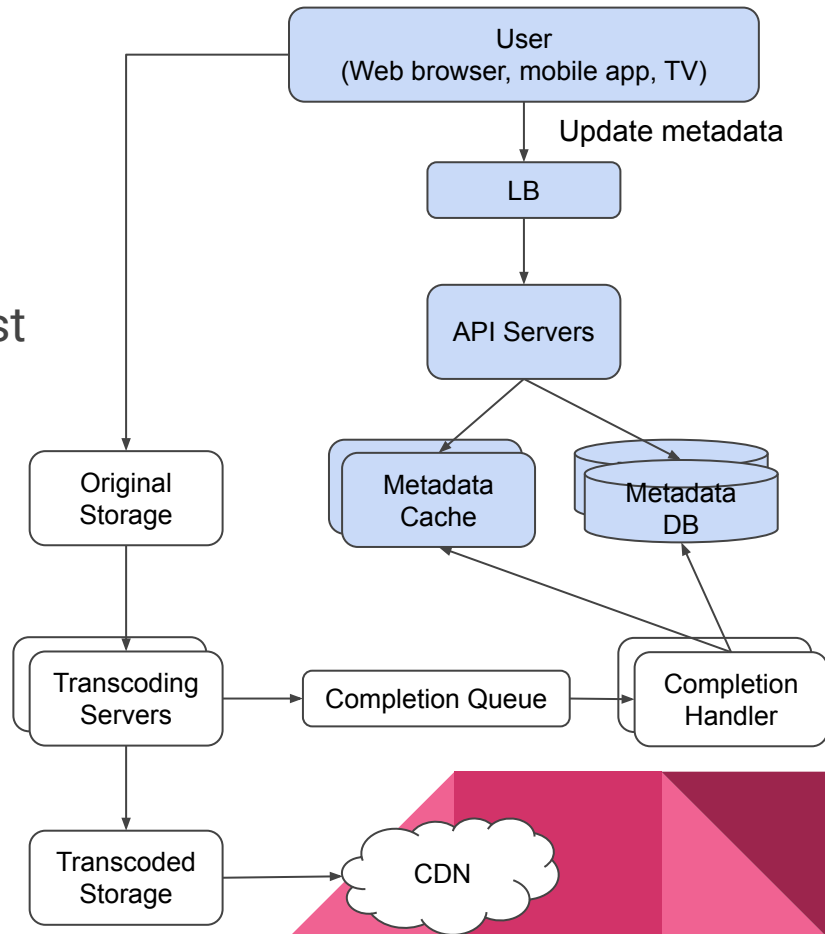
1	Videos: uploaded → original storage
2	Fetch videos from original storage & start transcoding
3	Transcoding complete → parallel processing followings
	3.a transcoded video → transcoded storage → CDN
	3.b completion event → queue
	3.b.1 Handler (workers): pull event data from queue - Update metadata (a. Database, b. cache)
4	API server → (video uploaded, ready for streaming) → client



### 3. High Level Design

(Video Uploading Workflow: update metadata)


- While a file is being uploaded to the original storage, the client sends a request to update the video metadata
- Request Body (Metadata)
  - File Name
  - Size
  - Format
  - Etc.
- API Servers
  - Update the metadata cache & DB



# 3. High Level Design

## (Video Streaming Workflow)

- **Purpose**
  - Whenever you watch a video on YouTube, it starts streaming immediately
  - Don't wait until whole video is downloaded on device.
- **Videos are streamed from CDN directly.**
  - The edge server closest to you deliver video.
  - Very little latency
- **Different streaming protocols support different video encodings and playback players.**
  - MPEG-DASH
  - Apple HLS (HTTP Live Streaming)
  - Microsoft Smooth Streaming
  - Adobe HDS (HTTP Dynamic Streaming)



## 4. Detailed Design

- High Level Design

- Video Uploading
- Video Streaming



- Design Deep Dive

- Refine both workflows w/
  - Optimizations
  - Error Handling Mechanisms

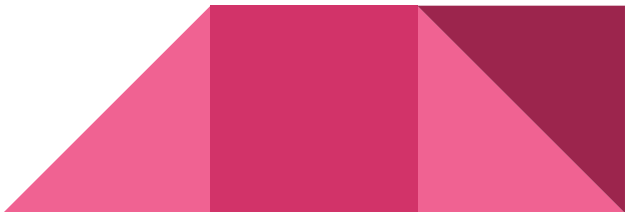
## 4. Detailed Design - Video Transcoding

### Reasons

- **Space** Optimization: HD raw video (60 frames/sec: a few hundreds GB)
- Device/Browser **Compatibility**: they only support certain types of video format
- **High-quality** videos: deliver high/low resolution video per network bandwidth
- **Smooth UX** (for playing continuously): switch quality automatically/manually

### Encoding Format

- **Container**
  - Basket for containing video file, audio, and metadata
  - Format by file extension (.AVI, .MOV, .MP4)
- **Codecs**
  - Compression/Decompression Algorithms: H.264, VP9, HEVC
  - To **reduce video size** while preserving video quality



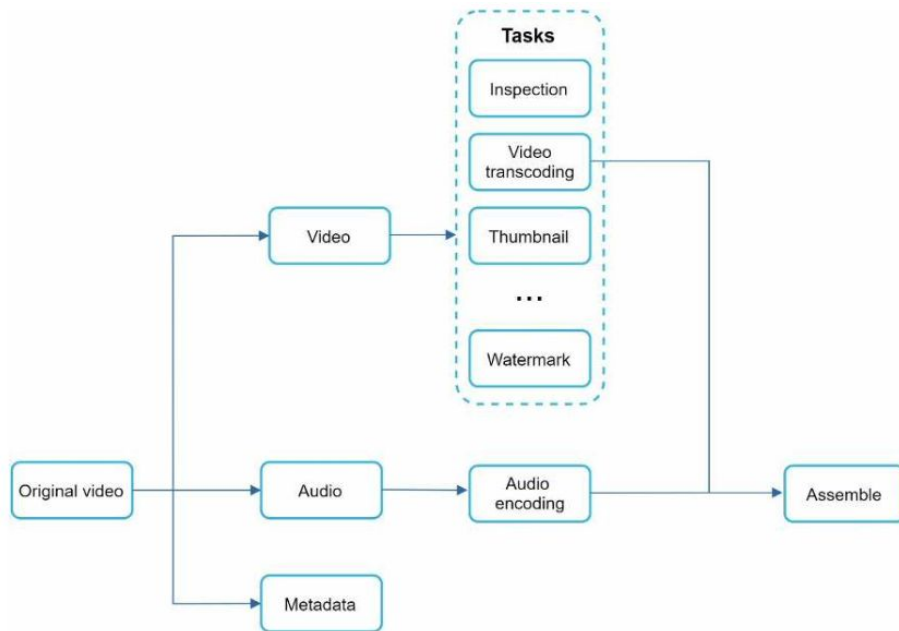
## 4. Detailed Design - DAG Model (Directed Acyclic Graph)

### Transcoding a video

- Computationally expensive
- Time-consuming
- Different content creators: different video processing requirements
  - Watermarks on top of their videos
  - Thumbnail images themselves
  - Someone upload HD videos, whereas others do not

### Directed acyclic graph (DAG) programming model

- Purpose
  - Support different video processing pipelines
  - Maintain high parallelism
- Add some level of abstraction
- Let client SWEs define what tasks to execute in stages
  - DAG can be executed sequentially or parallelly





## 4. Detailed Design - DAG Model

- **Inspection**

- Make sure the videos have good quality and aren't malformed

- **Video encodings**

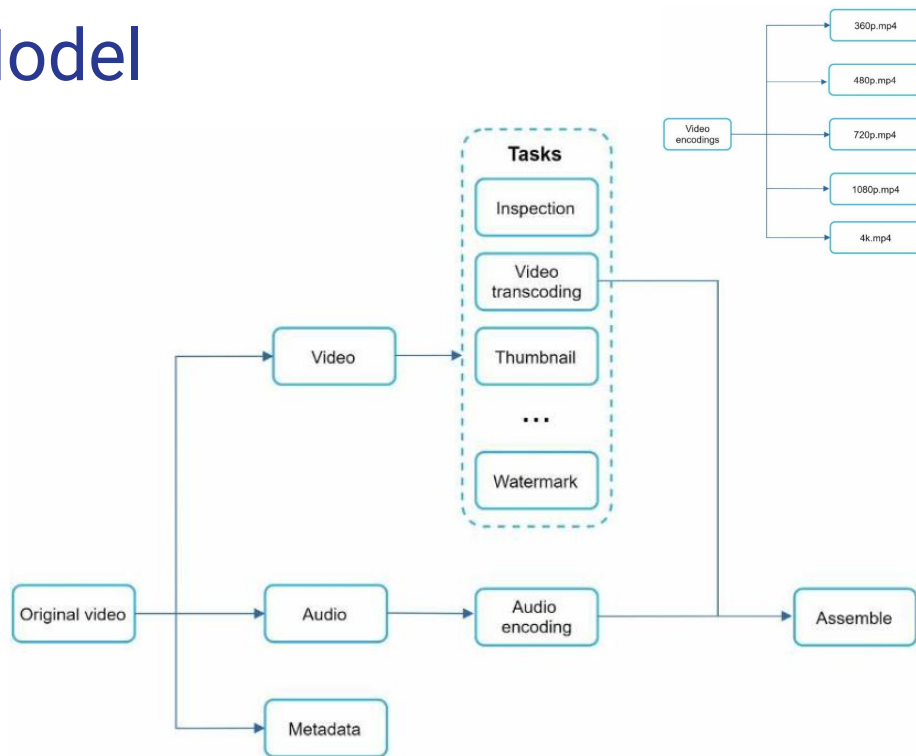
- Converted to support different resolutions, codec, bitrates, etc.

- **Thumbnail**

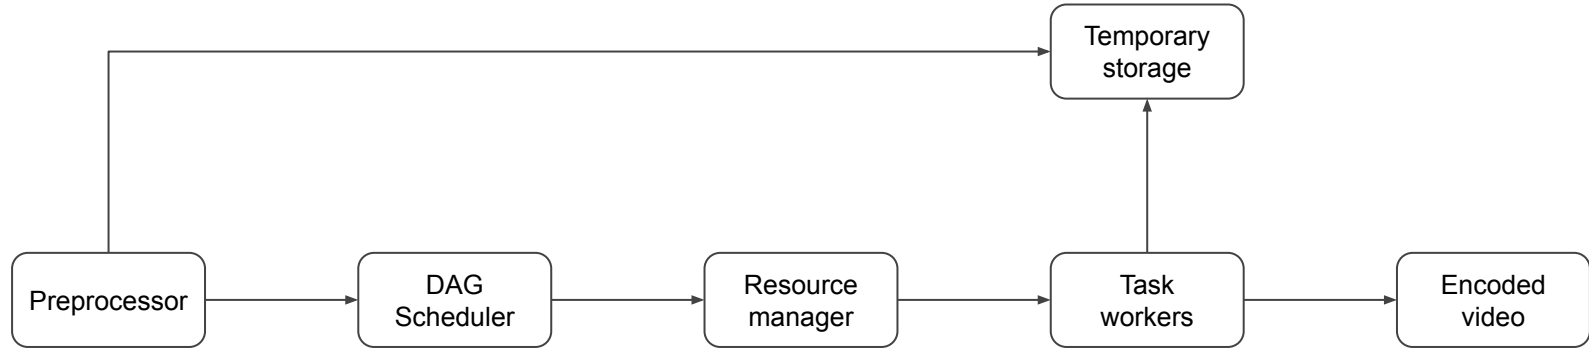
- Uploaded by user
- Automatically generated by system

- **Watermark**

- Image overlay on top of video (contains identifying info about your video)




## 4. Detailed Design - Video Transcoding Architecture



## 4. Detailed Design

(Video Transcoding Architecture)



- **Video Splitting**

- Groups of Pictures(GOP) Alignment: group/chunk of frames in a specific order
- Chunk: independently playable (a few secs)
- Not supported browser/old-devices: preprocessor split videos by GOP alignment

- **DAG Generation**

- Based on configuration files

- **Cache Data**

- Cache for segmented videos
- For better reliability: store GOPs and metadata in temporary storage.
  - If video encoding fails → use persisted data for retry operations



Figure 14-12


This DAG representation is generated from the two configuration files below (Figure 14-13):

```
task {
  name 'download-input'
  type 'Download'
  input {
    url config.url
  }
  output { it->
    context.inputVideo = it.file
  }
  next 'transcode'
}
```

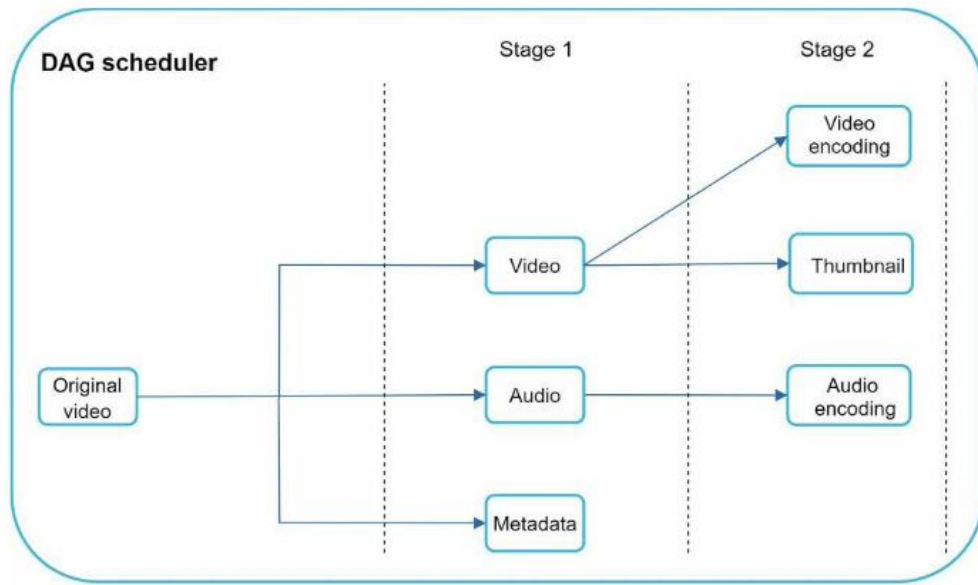
```
task {
  name 'transcode'
  type 'Transcode'
  input {
    input context.inputVideo
    config config.transConfig
  }
  output { it->
    context.file = it.outputVideo
  }
}
```

## 4. Detailed Design

(Video Transcoding Architecture)




- Split a DAG into tasks' stages
  - Original video → video, audio, metadata
  - Video file → 2 tasks (video encoding, thumbnail)
  - Audio file → audio encoding
- Put them into task queue in resource manager

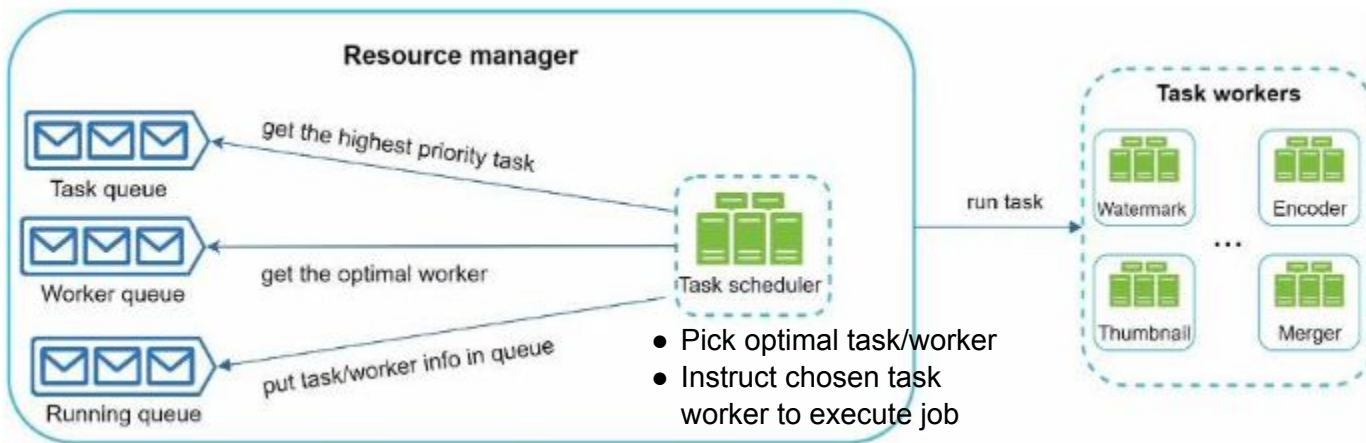


## 4. Detailed Design

(Video Transcoding Architecture)




Managing resource allocation efficiency



## 4. Detailed Design

(Video Transcoding Architecture)



### Task Workers


- Run tasks (defined in DAG)
- Different task workers may run different tasks

### Temporary Storage

- System **Choice**: depends on factors (type, size, access freq., life span)
- **Metadata**: Frequency: Highly accessed by workers; Size: small → In memory caching: good idea
- **Video/Audio Data**: Blob storage; Data is freed up once corresponding video processing is complete

### Encoded Video

- Final output of encoding pipeline: funny\_720p.mp4

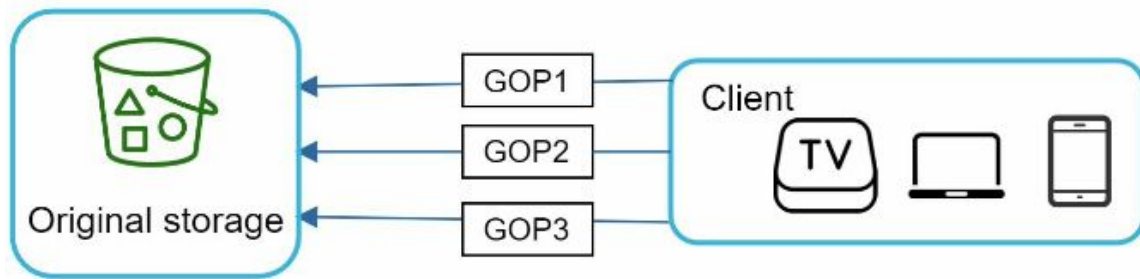


## 4. Detailed Design - Optimization

Upload a video (whole unit) → Inefficient → Split a video → smaller chunks



Previous upload failure → fast resumable uploads (Client to implement to improve upload speed)



### Speed Optimization

- **Parallelize video uploading**
  - Place upload centers close to users
  - Parallelism everywhere

### Safety Optimization

- Pre-signed upload URL
- Protect your videos

### Cost-Saving Optimization

## 4. Detailed Design - Optimization

Setting up multiple upload centers across the globe.

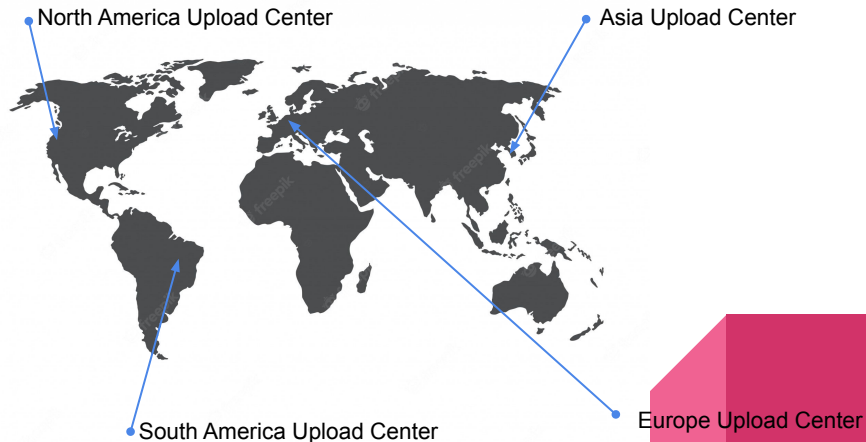
### Speed Optimization

- Parallelize video uploading
- **Place upload centers close to users**
- Parallelism everywhere

### Safety Optimization

- Pre-signed upload URL
- Protect your videos

### Cost-Saving Optimization





## 4. Detailed Design - Optimization

As-Is: Dependency: original storage → CDN: parallelism difficult

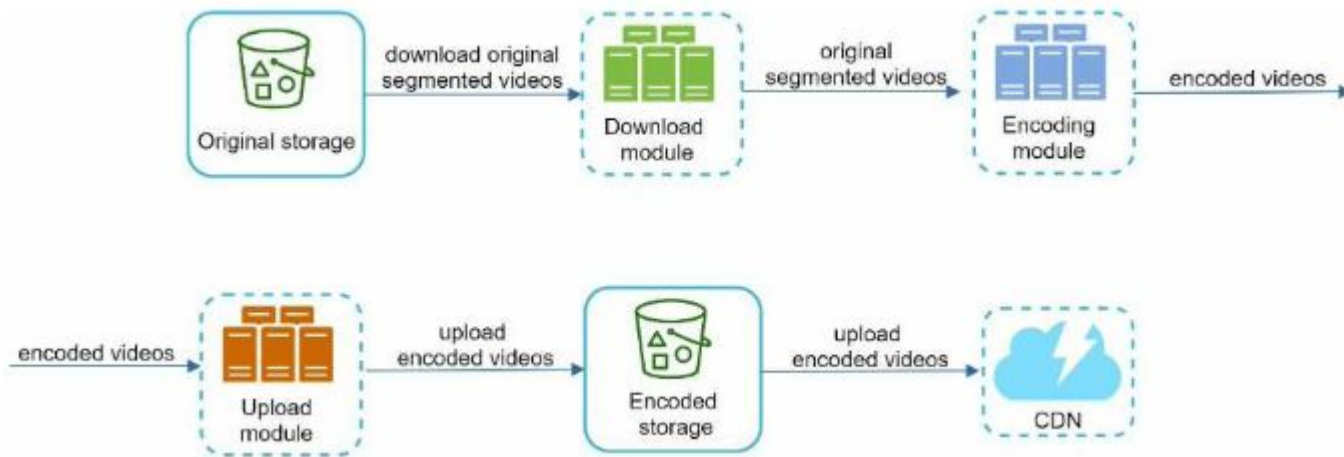
### Speed Optimization

- Parallelize video uploading
- Place upload centers close to users
- **Parallelism everywhere**

### Safety Optimization

- Pre-signed upload URL
- Protect your videos

### Cost-Saving Optimization



## 4. Detailed Design - Optimization

- Encoding module: wait before message queue is introduced
- After msg queue is introduced
  - Encoding module: no wait for output of Download module
  - Many events in the queue → Encoding module execute jobs in parallel

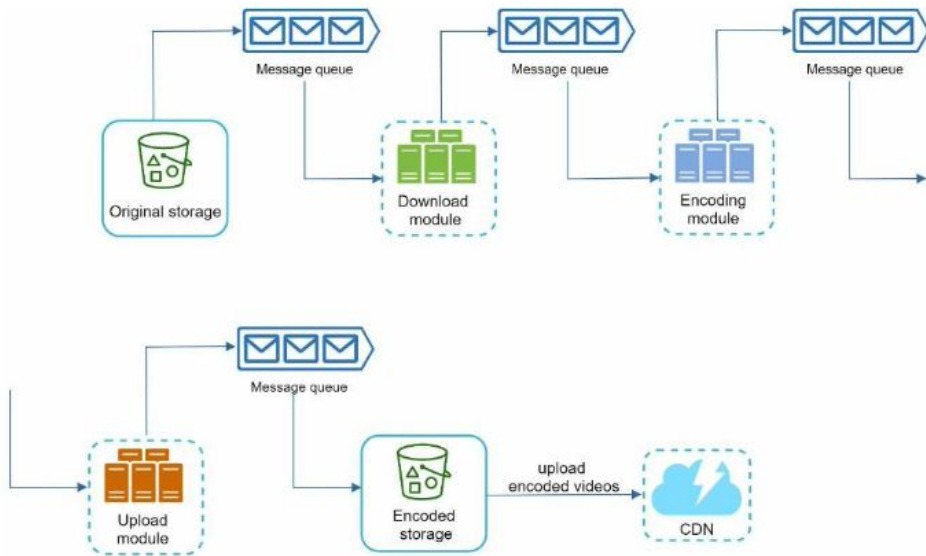
### Speed Optimization

- Parallelize video uploading
- Place upload centers close to users
- **Parallelism everywhere**

### Safety Optimization

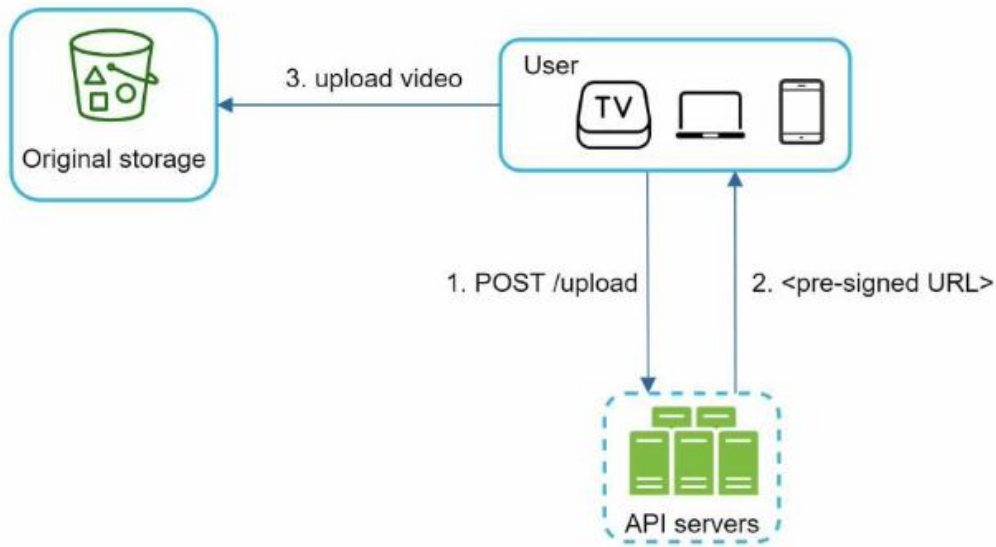
- Pre-signed upload URL
- Protect your videos

### Cost-Saving Optimization



## 4. Detailed Design - Optimization

Ensure only authorized users upload videos → right location.



### Speed Optimization

- Parallelize video uploading
- Place upload centers close to users
- Parallelism everywhere

### Safety Optimization

- **Pre-signed upload URL**
- Protect your videos

### Cost-Saving Optimization

- Amazon S3
- Azure blob storage (via Shared Access Signature)

## 4. Detailed Design - Optimization

### Content Makers

- Reluctant to upload to post videos online
- Reason: fear original videos → stolen

### To **Protect** Copyrighted Videos:

- Digital rights management (**DRM**) system: Apple FairPlay, Google Widevine, Microsoft PlayReady
- **AES** encryption
  - Encrypt a video, configure an authZ policy
  - Decrypt upon playback; Authorized users can only watch an encrypted video
- **Visual watermarking:**
  - Image on top of video (w/ identifying info for your video)
  - E.g., Company logo or name

#### Speed Optimization

- Parallelize video uploading
- Place upload centers close to users
- Parallelism everywhere

#### Safety Optimization

- Pre-signed upload URL
- **Protect your videos**

#### Cost-Saving Optimization

## 4. Detailed Design - Optimization

### CDN

- Pros: Crucial component for fast video deliver on global scale
- Cons: Expensive (size large?) → How to reduce cost?

**Optimization** based on content popularity, user access pattern, video size

- **CDN: Most popular videos**
  - Some videos: only popular in certain regions; no need to distribute them to other regions
  - Build your own CDN and partner w/ ISPs → improve viewing experience & reduce bandwidth charges
- **Video Servers: other less popular videos**
  - May not need to store many encoded video versions

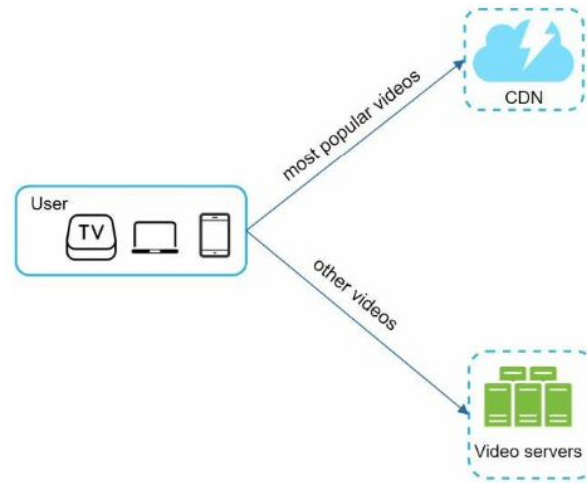
#### Speed Optimization

- Parallelize video uploading
- Place upload centers close to users
- Parallelism everywhere

#### Safety Optimization

- Pre-signed upload URL
- Protect your videos

#### Cost-Saving Optimization

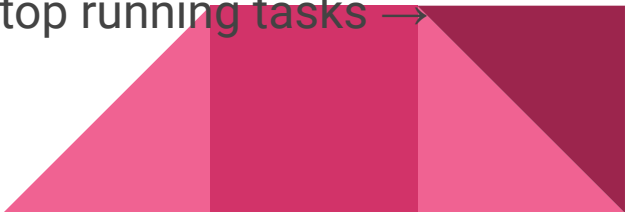


## 4. Detailed Design - Error Handling Mechanism

### Large-scale System Errors: Unavoidable

- Build highly fault-tolerant system
- Handle errors gracefully
- Recover fast

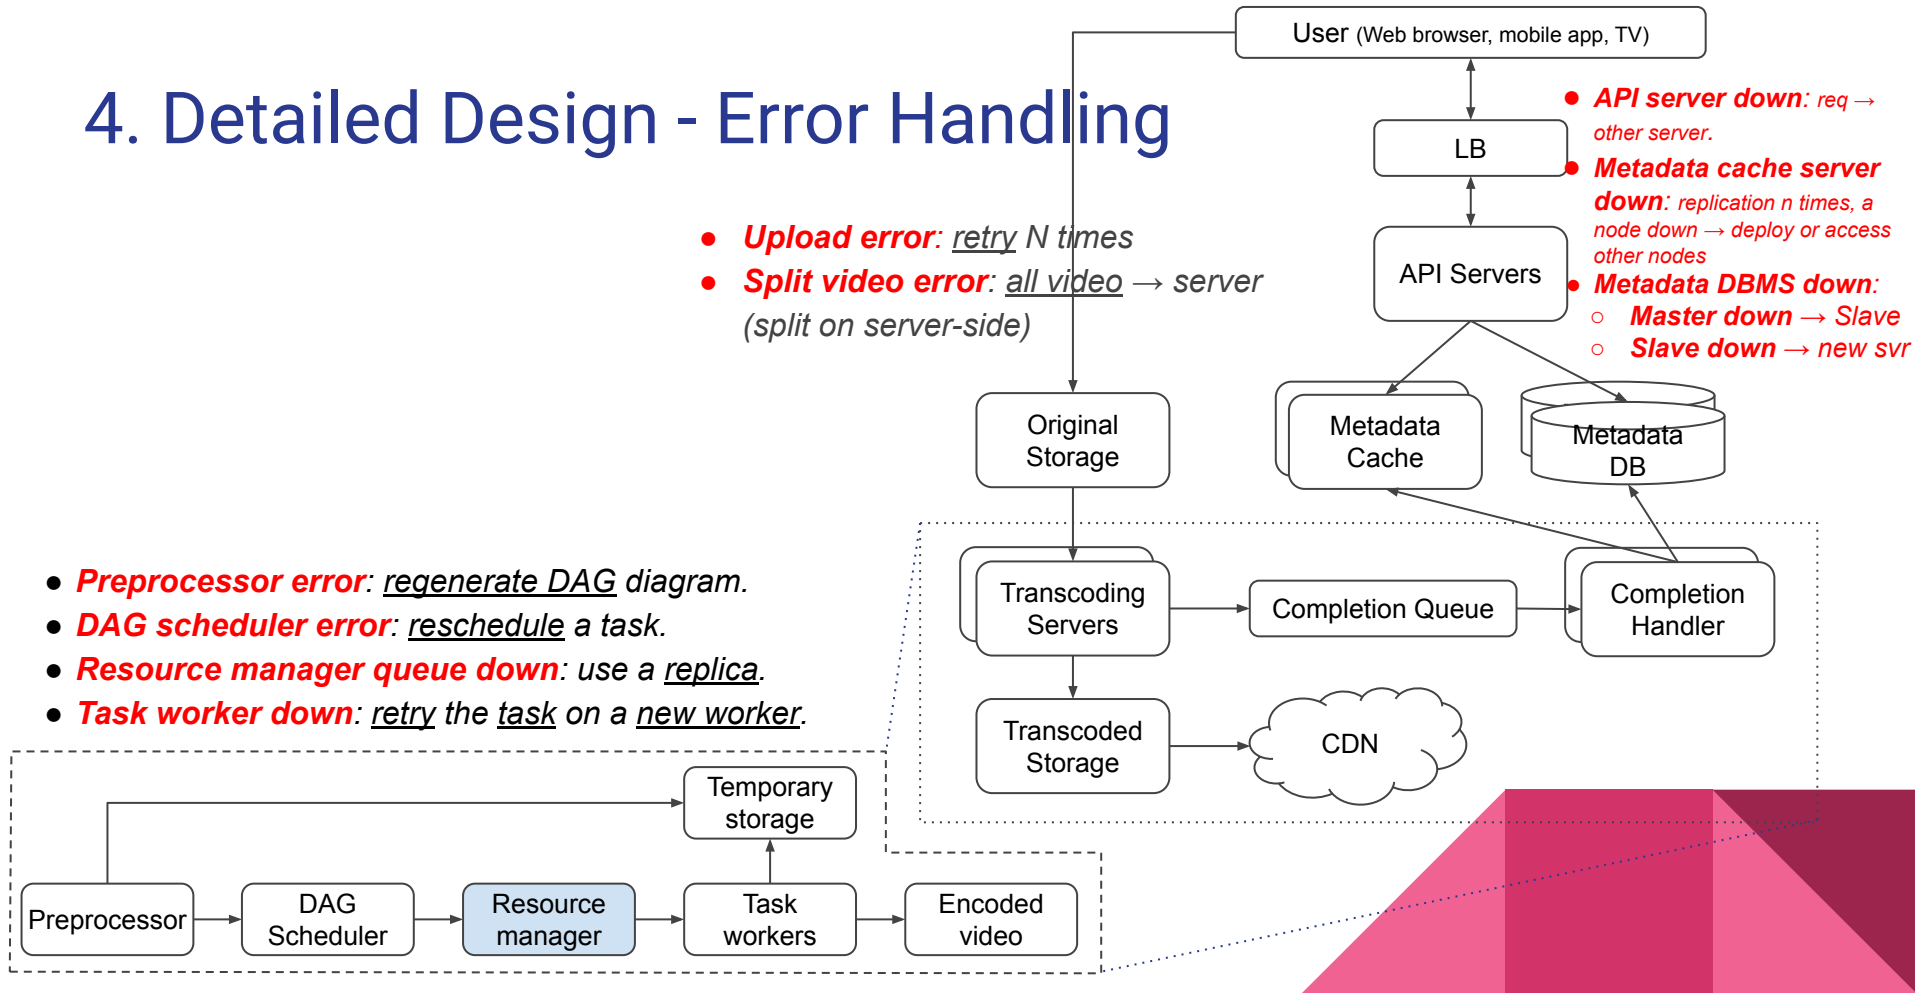
### Error Types

- **Recoverable Error:** video segment fails to transcode → retry n times → fail → error code to client
  - **Non-recoverable Error:** malformed video format → stop running tasks → error code to client
- 

## 4. Detailed Design - Error Handling

- **Upload error:** retry  $N$  times
- **Split video error:** all video → server (split on server-side)

- **Preprocessor error:** regenerate DAG diagram.
- **DAG scheduler error:** reschedule a task.
- **Resource manager queue down:** use a replica.
- **Task worker down:** retry the task on a new worker.




# 5. Wrap Up

## What We Did

- Architecture design for video streaming services (like YouTube)

## Additional Points

- **Scale API:** API servers (stateless) → easy to scale horizontally
- **Scale DB:** replication, sharding
- **Live Streaming:**
  - Higher latency → need different streaming protocol
  - Lower requirement for parallelism → small chunks (already processed in real-time)
  - Different sets of error handling: too much time (not acceptable)
- **Video Takedowns:**
  - Copyrights, pornography, other illegal acts: shall be removed
  - Some can be discovered during upload process or via user flagging





# References

- Alex Xu, System Design Interview (Insider's Guide)
- [Scaling Patterns for Netflix's Edge](#)

